Socket Programming

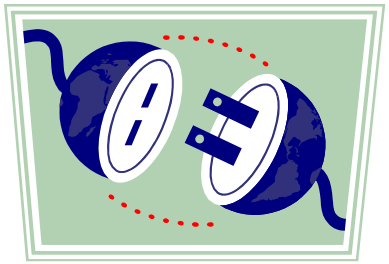

Rohan Murty

Hitesh Ballani

Last Modified:

2/8/2004 8:31:51 AM

Slides adapted from Prof. Matthews' slides from 2003SP

Socket programming

Goal: learn how to build client/server application that communicate using sockets

Socket API

- introduced in BSD4.1 UNIX, 1981
- Sockets are explicitly created, used, released by applications
- client/server paradigm
- two types of transport service via socket API:
  - unreliable datagram
  - reliable, byte stream-oriented

socket

a *host-local*, *application-created/owned*, *OS-controlled* interface (a "door") into which application process can both send and receive messages to/from another (remote or local) application process

Sockets

Socket: a door between application process and end-end-transport protocol (UCP or TCP)

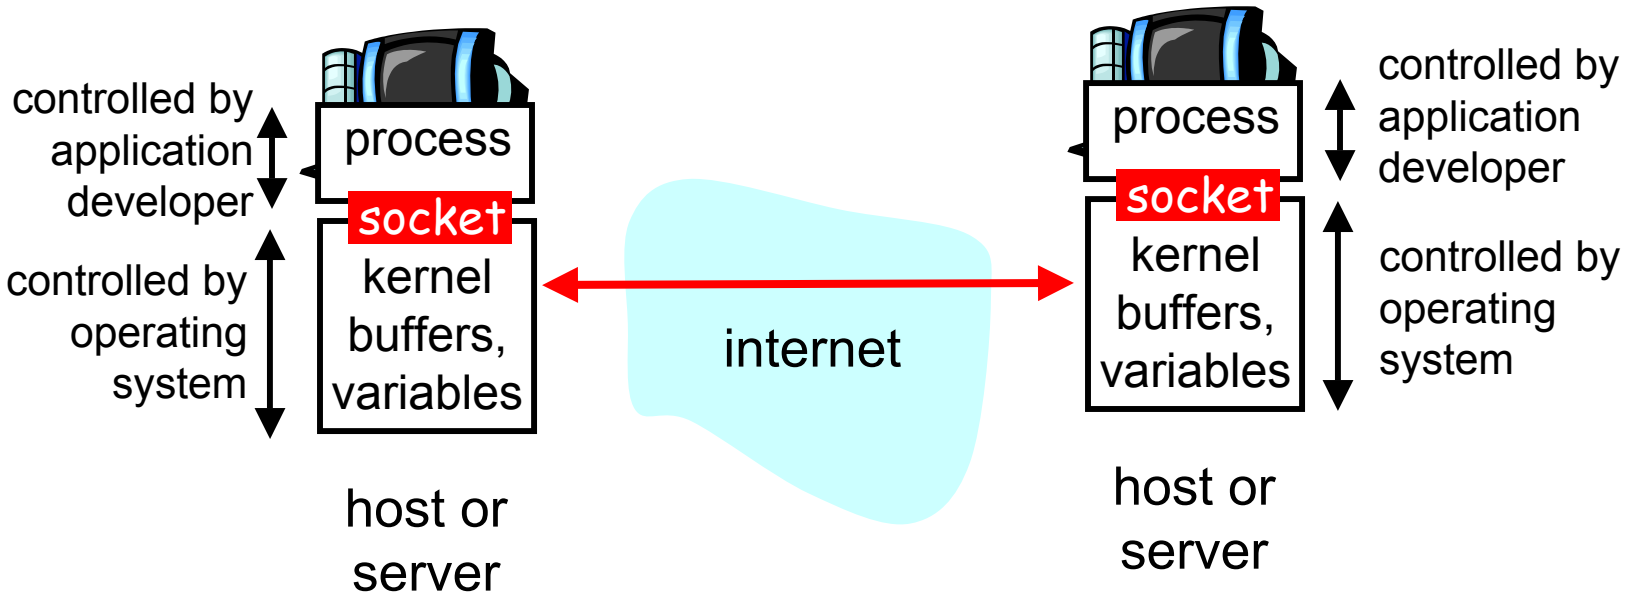

Languages and Platforms

Socket API is available for many languages on many platforms:

- C, Java, Perl, Python,…
- *nix, Windows,…

Socket Programs written in any language and running on any platform can communicate with each other!

Writing communicating programs in different languages is a good exercise

Decisions

- Before you go to write socket code, decide
  - Do you want a TCP-style reliable, full duplex, connection oriented channel? Or do you want a UDP-style, unreliable, message oriented channel?
  - Will the code you are writing be the client or the server?
    - Client: you assume that there is a process already running on another machines that you need to connect to.
    - Server: you will just start up and wait to be contacted

Socket programming with TCP

Client must contact server

- server process must first be running
- server must have created socket (door) that welcomes client's contact

Client contacts server by:

- creating client-local TCP socket
- specifying IP address, port number of server process

- When client creates socket: client TCP establishes connection to server TCP
- When contacted by client, server TCP creates new socket for server process to communicate with client
  - Frees up incoming port
  - allows server to talk with multiple clients

application viewpoint

TCP provides reliable, in-order transfer of bytes ("pipe") between client and server

Pseudo code TCP client

- Create socket, connectSocket
- Do an active connect specifying the IP address and port number of server
- Read and Write Data Into connectSocket to Communicate with server
- Close connectSocket

Pseudo code TCP server

- Create socket (serverSocket)
- Bind socket to a specific port where clients can contact you
- Register with the kernel your willingness to listen that on socket for client to contact you
- Loop

    Accept new connection (connectSocket)

    Read and Write Data Into connectSocket to Communicate with client

    Close connectSocket

  End Loop
- Close serverSocket

Example: Java client (TCP)

```
import java.io.*;
import java.net.*;
class TCPClient {

    public static void main(String argv[]) throws Exception
    {
        String sentence;
        String modifiedSentence;

        BufferedReader inFromUser =
          new BufferedReader(new InputStreamReader(System.in));

        Socket clientSocket = new Socket("hostname", 6789);

        DataOutputStream outToServer =
          new DataOutputStream(clientSocket.getOutputStream());
```

Create input stream

Create client socket, connect to server

Create output stream attached to socket

Example: Java client (TCP), cont.

Create input stream attached to socket

```
BufferedReader inFromServer =
  new BufferedReader(new
  InputStreamReader(clientSocket.getInputStream()));

sentence = inFromUser.readLine();

outToServer.writeBytes(sentence + '\n');

modifiedSentence = inFromServer.readLine();

System.out.println("FROM SERVER: " + modifiedSentence);

clientSocket.close();

    }
}
```

Send line to server

Read line from server

Example: Java server (TCP)

```
import java.io.*;
import java.net.*;

class TCPServer {

  public static void main(String argv[]) throws Exception
   {
     String clientSentence;
     String capitalizedSentence;

     ServerSocket welcomeSocket = new ServerSocket(6789);

     while(true) {

       Socket connectionSocket = welcomeSocket.accept();

      BufferedReader inFromClient =
        new BufferedReader(new
        InputStreamReader(connectionSocket.getInputStream()));
```

Create welcoming socket at port 6789

Wait, on welcoming socket for contact by client

Create input stream, attached to socket

Example: Java server (TCP), cont

Create output stream, attached to socket

```
DataOutputStream  outToClient =
  new DataOutputStream(connectionSocket.getOutputStream());
```

Read in line from socket

```
clientSentence = inFromClient.readLine();

capitalizedSentence = clientSentence.toUpperCase() + '\n';
```

Write out line to socket

```
outToClient.writeBytes(capitalizedSentence);
        }
    }
}
```

End of while loop, loop back and wait for another client connection

Client/server socket interaction: TCP (Java)

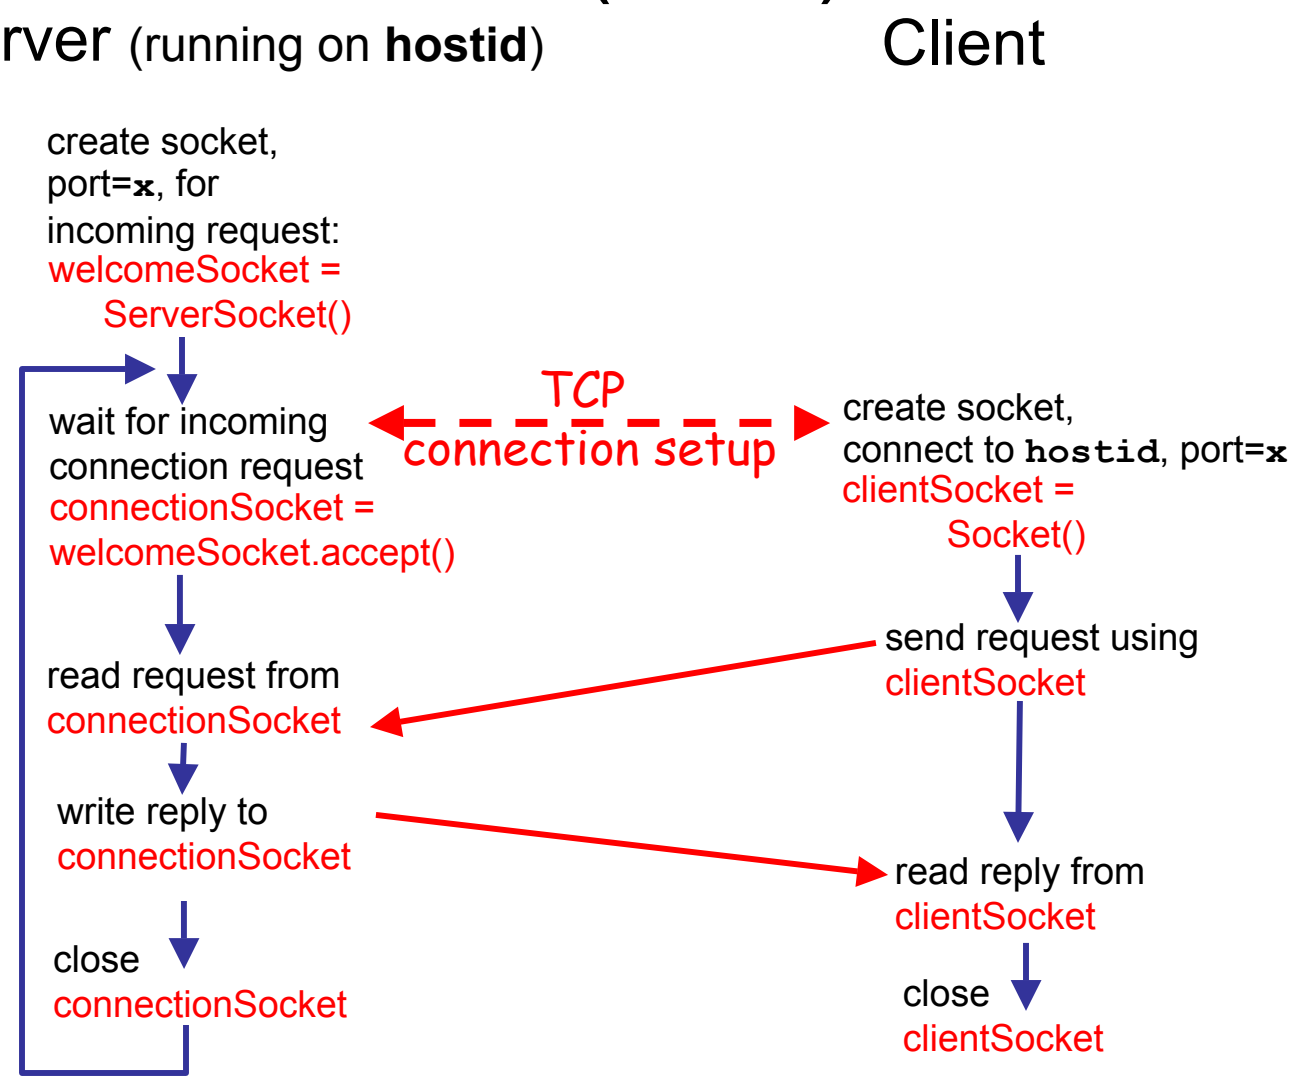

Queues

- We just saw a simple example, with one socket on the server handling incoming connections
- While the server socket is busy, incoming connections are stored in a queue until it can accept them
- Most systems maintain a queue length between 5 and 50
- Once the queue fills up, further incoming connections are refused until space in the queue opens up
- This is a problem in a situation where our server has to handle many concurrent incoming connections. Example: HTTP servers
  - Solution? Use concurrency

Concurrent TCP Servers

- Benefit comes in ability to hand off processing to another process
  - Parent process creates the “door bell” or “welcome” socket on well-known port and waits for clients to request connection
  - When a client does connect, fork off a child process to handle that connection so that parent process can return to waiting for connections as soon as possible
- Multithreaded server: same idea, just spawn off another thread rather than a full process
  - Threadpools?

Threadpools

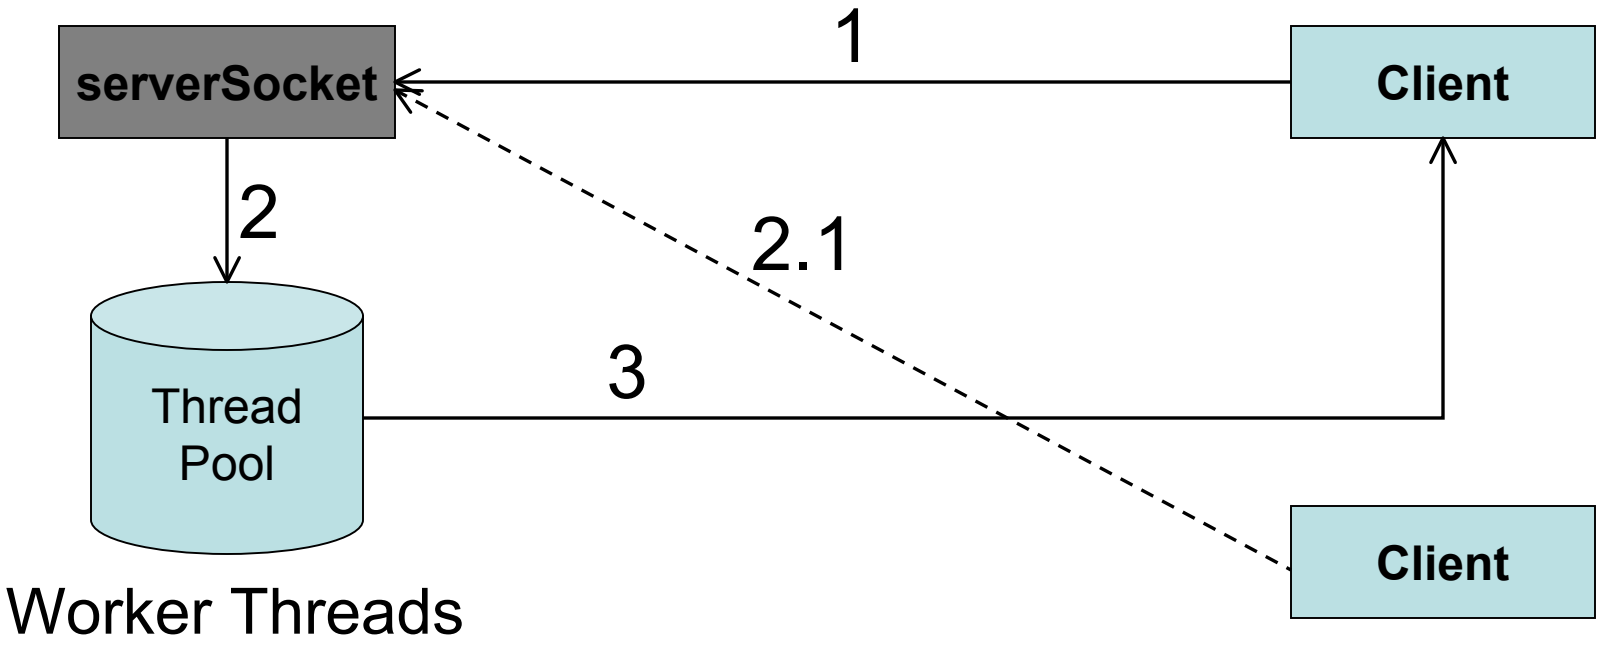

Socket programming with UDP

UDP: very different mindset than TCP

- no connection just independent messages sent
- no handshaking
- sender explicitly attaches IP address and port of destination
- server must extract IP address, port of sender from received datagram to know who to respond to

UDP: transmitted data may be received out of order, or lost

application viewpoint

UDP provides unreliable transfer of groups of bytes ("datagrams") between client and server

Pseudo code UDP server

- Create socket
- Bind socket to a specific port where clients can contact you
- Loop

  (Receive UDP Message from client x)+

  (Send UDP Reply to client x)*
- Close Socket

Pseudo code UDP client

- Create socket
- Loop

  (Send Message To Well-known port of server)+

  (Receive Message From Server)
- Close Socket

Example: Java client (UDP)

```
import java.io.*;
import java.net.*;

class UDPClient {
    public static void main(String args[]) throws Exception
    {

      BufferedReader inFromUser =
        new BufferedReader(new InputStreamReader(System.in));

      DatagramSocket clientSocket = new DatagramSocket();

      InetAddress IPAddress = InetAddress.getByName("hostname");

      byte[] sendData = new byte[1024];
      byte[] receiveData = new byte[1024];

      String sentence = inFromUser.readLine();

      sendData = sentence.getBytes();
```

Create input stream

Create client socket

Translate hostname to IP address using DNS

Example: Java client (UDP), cont.

Create datagram with data-to-send, length, IP addr, port

Send datagram to server

Read datagram from server

```
                DatagramPacket sendPacket =
                  new DatagramPacket(sendData, sendData.length, IPAddress, 9876);

                clientSocket.send(sendPacket);

                DatagramPacket receivePacket =
                   new DatagramPacket(receiveData, receiveData.length);

                clientSocket.receive(receivePacket);

                String modifiedSentence =
                    new String(receivePacket.getData());

                System.out.println("FROM SERVER:" + modifiedSentence);
                clientSocket.close();
                }

            }
```

Example: Java server (UDP)

```
import java.io.*;
import java.net.*;

class UDPServer {
  public static void main(String args[]) throws Exception
    {

      DatagramSocket serverSocket = new DatagramSocket(9876);

      byte[] receiveData = new byte[1024];
      byte[] sendData  = new byte[1024];

      while(true)
        {

        DatagramPacket receivePacket =
           new DatagramPacket(receiveData, receiveData.length);

        serverSocket.receive(receivePacket);
```

Create datagram socket at port 9876

Create space for received datagram

Receive datagram

Example: Java server (UDP), cont

```
String sentence = new String(receivePacket.getData());

InetAddress IPAddress = receivePacket.getAddress();

int port = receivePacket.getPort();

                String capitalizedSentence = sentence.toUpperCase();

sendData = capitalizedSentence.getBytes();

DatagramPacket sendPacket =
   new DatagramPacket(sendData, sendData.length, IPAddress,
                      port);

serverSocket.send(sendPacket);
    }
  }

}
```

Get IP addr port #, of sender

Create datagram to send to client

Write out datagram to socket

End of while loop, loop back and wait for another datagram

Client/server socket interaction: UDP

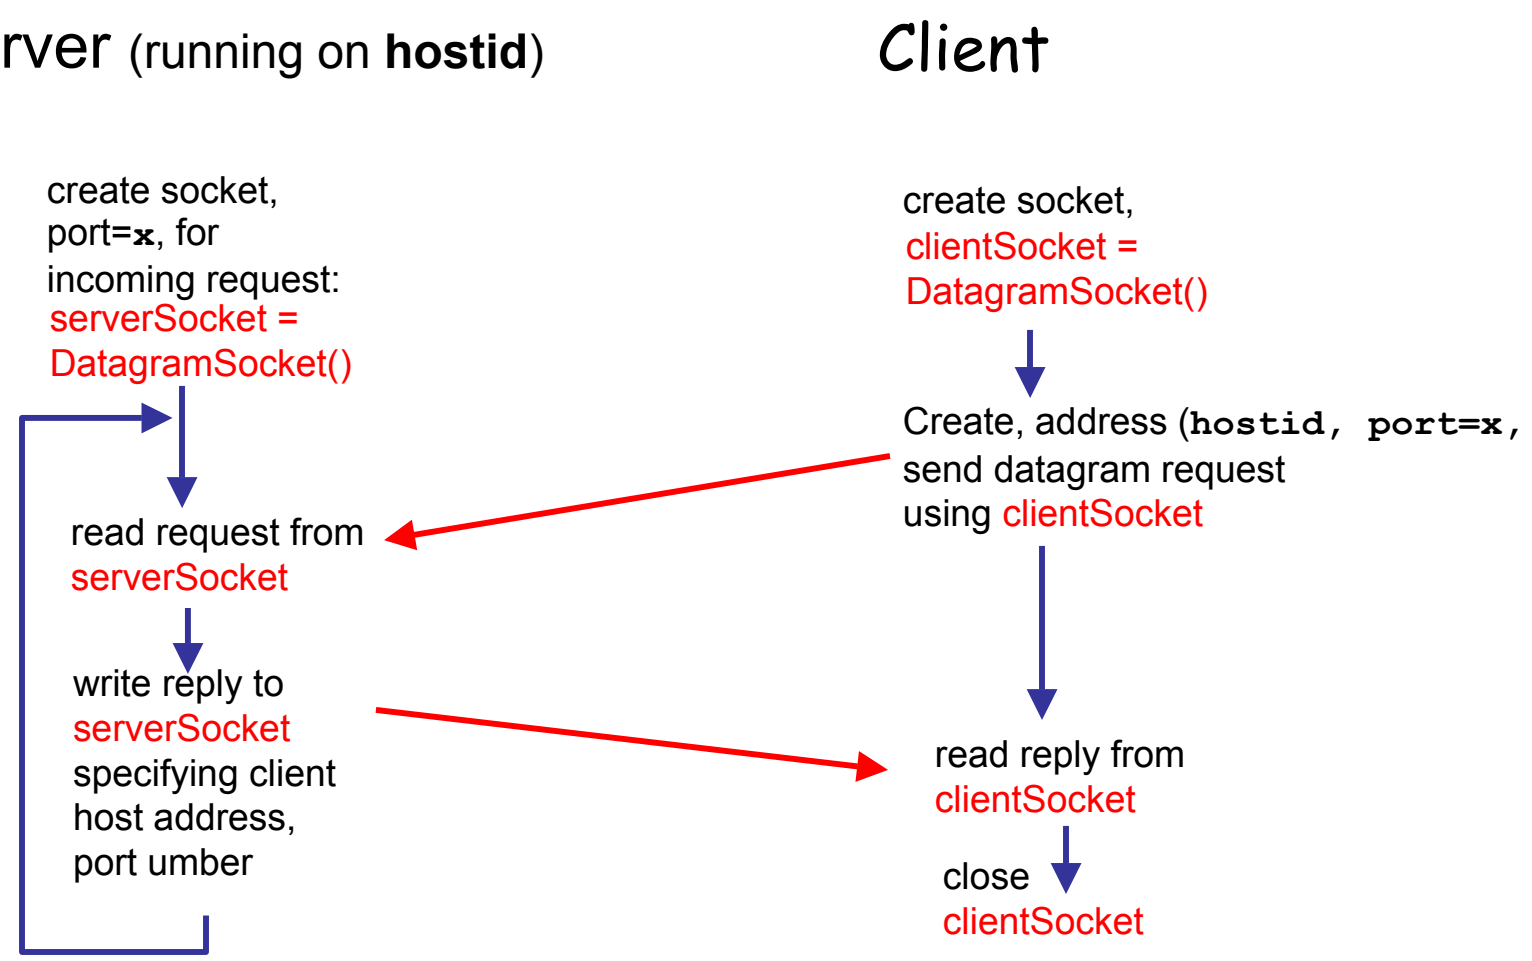

UDP Server vs Client

- Server has a well-known port number
- Client initiates contact with the server
- Less difference between server and client code than in TCP
  - Both client and server bind to a UDP socket
  - Not accept for server and connect for client
- Client send to the well-known server port; server extracts the client's address from the datagram it receives

TCP vs UDP

- TCP can use read/write (or recv/send) and source and destination are implied by the connection; UDP must specify destination for each datagram
  - Sendto, recvfrm include address of other party
- TCP server and client code look quite different; UDP server and client code vary mostly in who sends first

Socket Types

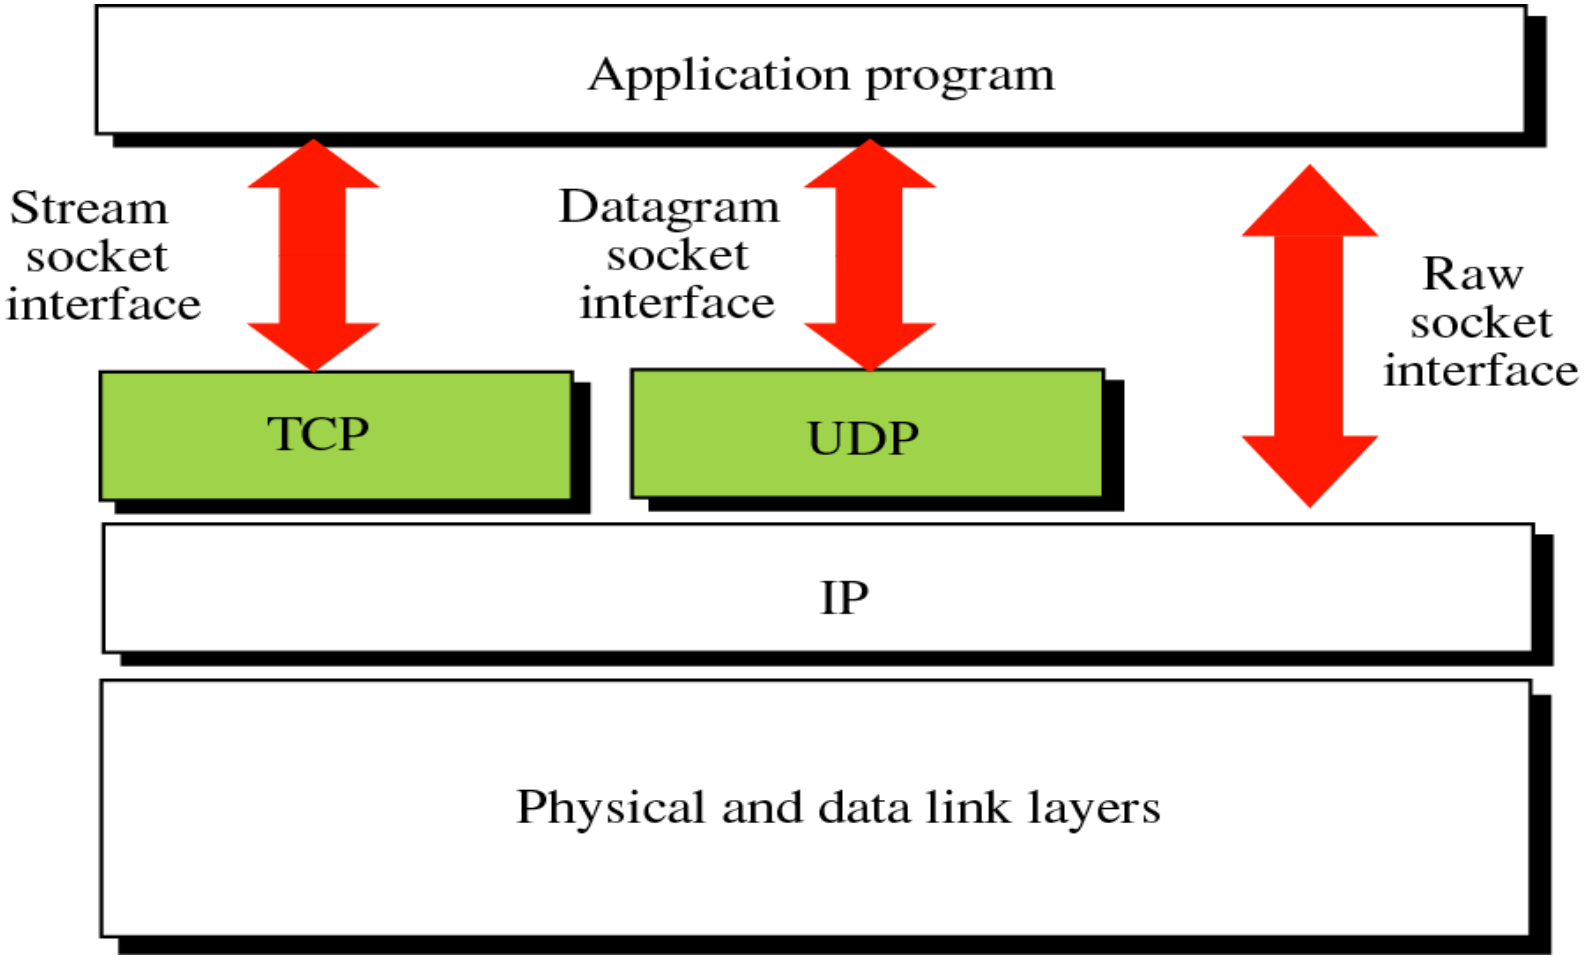

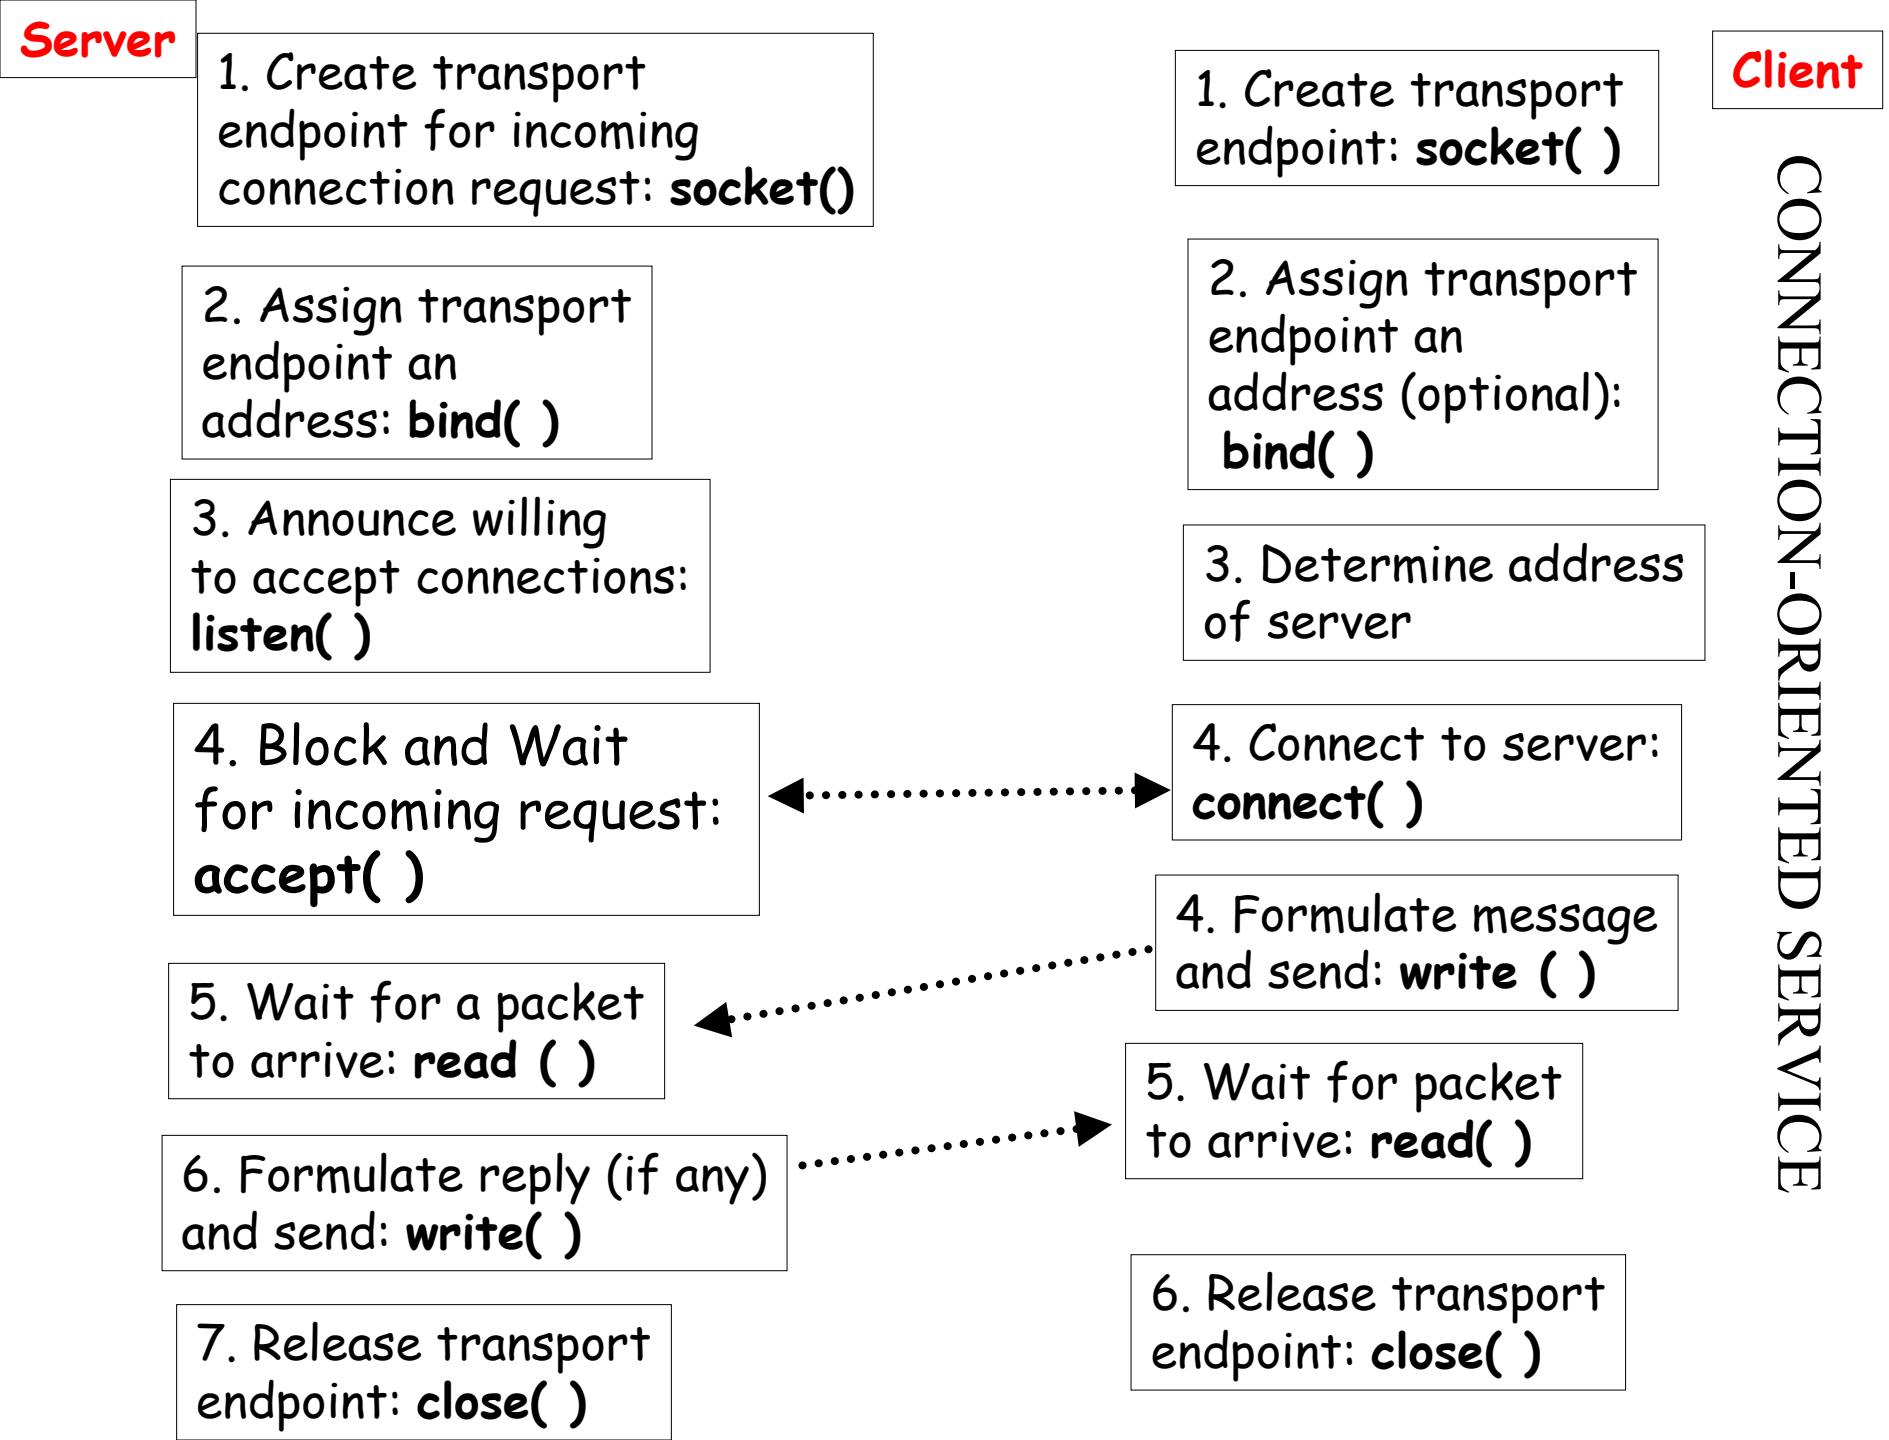
Server
1. Create transport endpoint for incoming connection request: socket()
2. Assign transport endpoint an address: bind( )
3. Announce willing to accept connections: listen( )
4. Block and Wait for incoming request: accept( )
5. Wait for a packet to arrive: read ( )
6. Formulate reply (if any) and send: write( )
7. Release transport endpoint: close( )
Client
1. Create transport endpoint: socket( )
2. Assign transport endpoint an address (optional): bind( )
3. Determine address of server
4. Connect to server: connect( )
4. Formulate message and send: write ( )
5. Wait for packet to arrive: read( )
6. Release transport endpoint: close( )
CONNECTION-ORIENTED SERVICE

Connectionless Service (UDP)

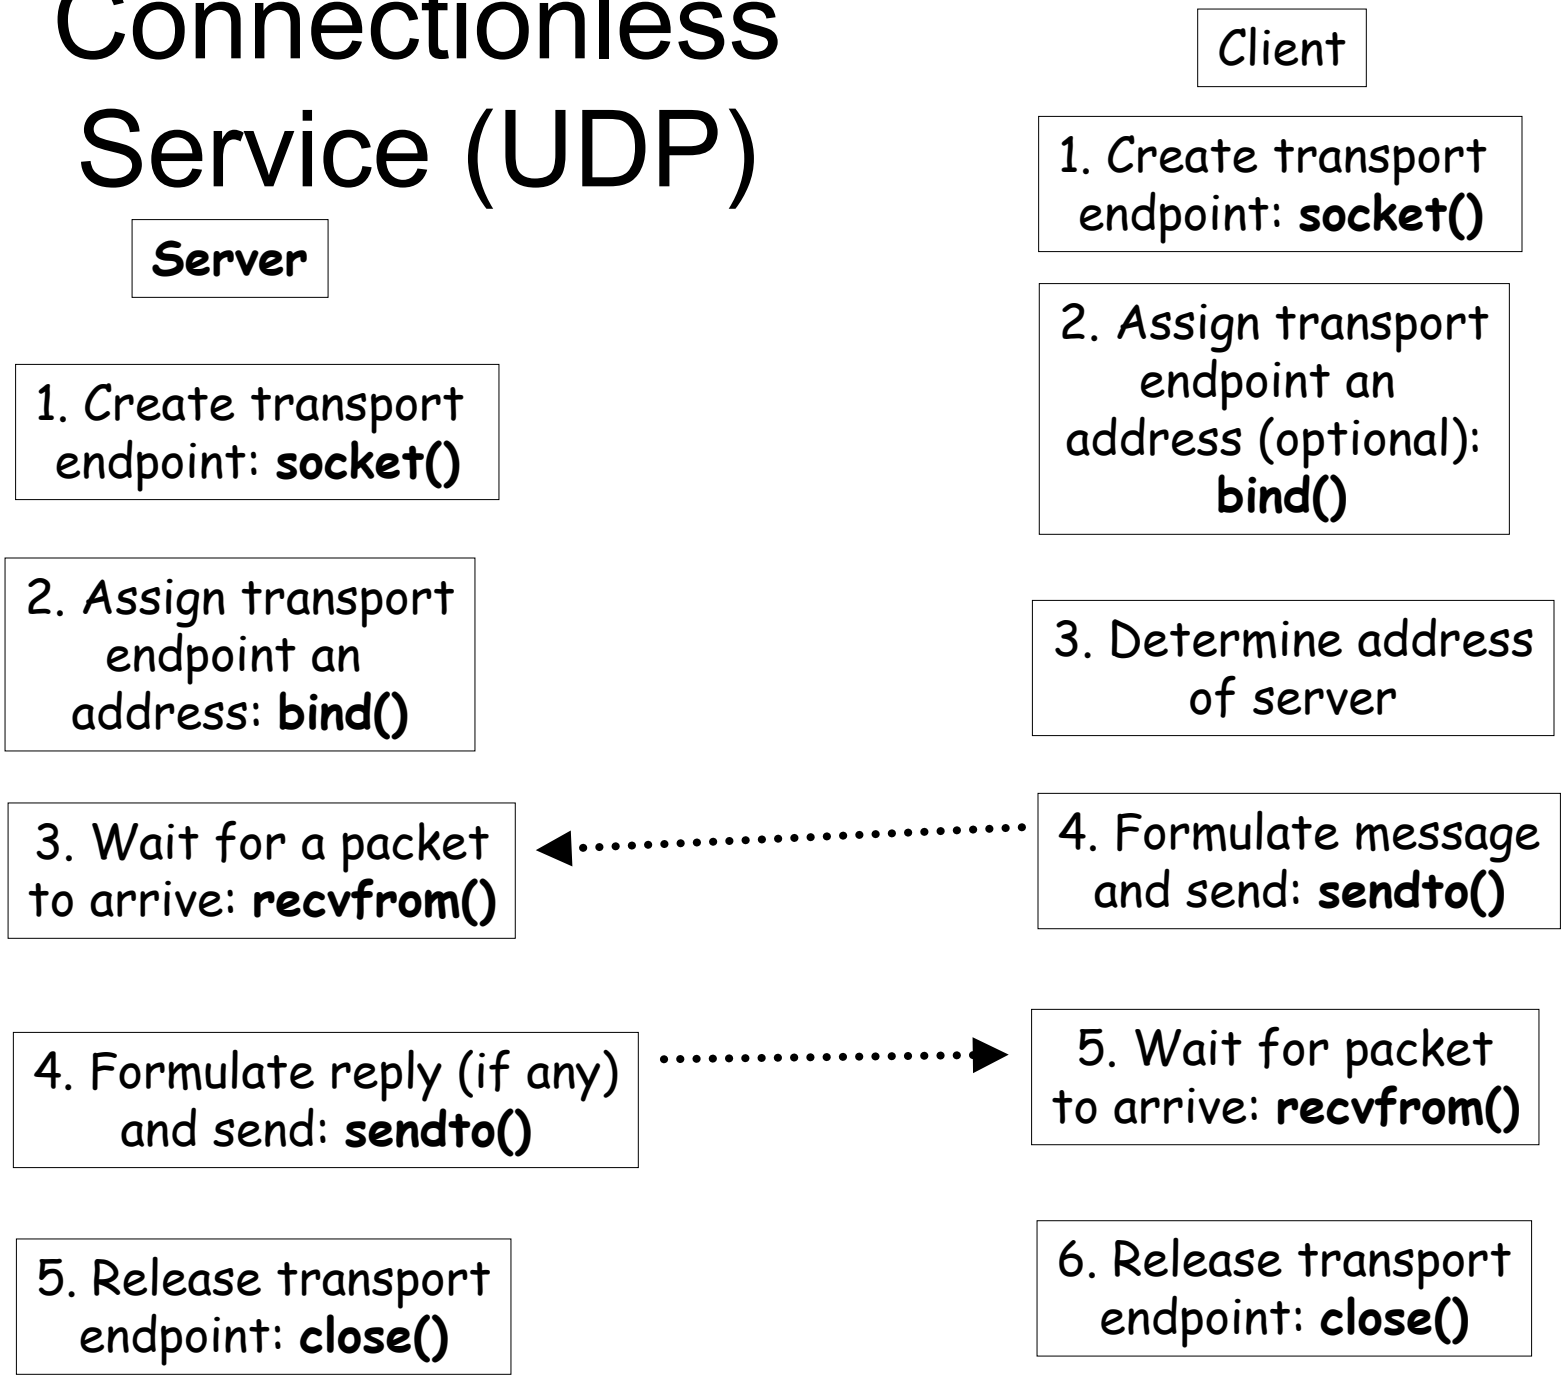

Non-blocking I/O

- By default, `accept`(), `recv`(), etc block until there's input
- What if you want to do something else while you're waiting?
- We can set a socket to not block (i.e. if there's no input an error will be returned)
- … or, we can tell the kernel to let us know when a socket is ready, and deal with it only then

PROJECT 1 : BASIC SOCKETS

AIM: Write a program (referred to as the **IP box**) that opens three sockets, two TCP and one UDP

2 TCP SOCKETS :

1. A **receive-config** socket : IP BOX acts as a Server (must be bound to a port you have to find, and the interface IP address)

2. A **send-config** socket : IP BOX acts as client.

(CONT …)

- 1 UDP SOCKETS

- App -- acts as the interface between the IP layer and the application

- It must be bound to an used port and the interface address

IP BOX OPERATION

- Send-config sockets connects to the Test Box and sends a “ready-to-test” command
- The Test Box then connects to recv-config socket and send a ‘\n’ terminated command which must be echoed
- The Test Box then sends UDP packets to app socket which must be echoed (Note : If the Test Box does not receive your echo, it retransmits the packet)

(cont …)

- On receiving both the echoes, the Test Box sends a “send-stat” command to the send-config socket
- The IP box sends a “list-of-stats”
- The Test Box then sends an exit message.